

R data.table Tutorial

(This tutorial is adapted from a technical presentation by Brian Silva, a data scientist at Uptake, LLC, a Chicago-based analytics firm focusing on heavy industry)

This tutorial presents **data.table**, an R package by Matt Dowle, which is an extension on R's default **data.frame** used for fast aggregation of large data (e.g. 100GB in RAM). We will compare its speed and syntax with that of **data.frame**.

Installation

First, make sure the **data.table** package is installed by running:

```
install.packages('data.table')
```

Data importing: fread vs. read.csv

The **fread** function, which has almost identical arguments as those of **read.csv**, works many times faster than **read.csv**. For example, a large (10GB, 100M rows x 10 cols) data set will take **read.csv** hours to read, while it will take **fread** only about 4-5 minutes.

By default **fread** will read a .csv file into a **data.table**, but you can also read it in as a **data.frame**.

```
dt <- fread('dataset.csv')           # returns data.table
df <- fread('dataset.csv', data.table=F) # returns data.frame
```

Syntax: data.table vs. data.frame

Let's start by making a large amount of data in both **data.frame** and **data.table** formats:

```
set.seed(35753)

numRows <- 10000000

# Create data with assets and dateTimes
assets <- paste('Asset', 1:1000, sep='_')

dateTimes <- seq.POSIXt(from=as.POSIXct('2014-01-01 00:00:00'),
                        to=as.POSIXct('2015-01-01 00:00:00'),
                        length.out=50000)

# Randomly select assets dateTimes and generate signals
DF <- data.frame(name=1:numRows,
                 assetId=sample(assets, numRows, replace=T),
                 dateTime=sample(dateTimes, numRows, replace=T),
                 signal1=runif(numRows),
                 signal2=rexp(numRows),
                 signal3=sample(c('ON', 'OFF'), numRows, replace=T),
                 stringsAsFactors=F)

DT <- as.data.table(DF)
```

Viewing the data

First of all, `data.table` inherits from `data.frame`, which means that it can be passed to any package that only accepts `data.frame`.

```
class(DT)
```

```
## [1] "data.table" "data.frame"
```

Let's take a look at the `data.frame` and `data.table` we created.

```
system.time(h <- head(DF))  # see how much time it takes to present head of data.frame
```

```
##      user  system elapsed  
##      0.00    0.02    0.02
```

```
h
```

```
##   name  assetId      dateTime  signal1  signal2 signal3  
## 1:    1 Asset_821 2014-08-22 22:52:14 0.73100637 0.7845882    OFF  
## 2:    2 Asset_19 2014-09-18 04:22:36 0.08924429 0.2062187    OFF  
## 3:    3 Asset_743 2014-09-24 18:24:37 0.56470635 0.9394910     ON  
## 4:    4 Asset_969 2014-06-26 14:07:17 0.14074313 0.1588975     ON  
## 5:    5 Asset_431 2014-10-06 02:19:34 0.25558412 0.3531577     ON  
## 6:    6 Asset_693 2014-02-23 22:45:14 0.24515170 0.1918058     ON
```

```
system.time(h <- head(DT))  # see how much time it takes to present head of data.table
```

```
##      user  system elapsed  
##       0       0       0
```

```
h
```

```
##   name  assetId      dateTime  signal1  signal2 signal3  
## 1:    1 Asset_821 2014-08-22 22:52:14 0.73100637 0.7845882    OFF  
## 2:    2 Asset_19 2014-09-18 04:22:36 0.08924429 0.2062187    OFF  
## 3:    3 Asset_743 2014-09-24 18:24:37 0.56470635 0.9394910     ON  
## 4:    4 Asset_969 2014-06-26 14:07:17 0.14074313 0.1588975     ON  
## 5:    5 Asset_431 2014-10-06 02:19:34 0.25558412 0.3531577     ON  
## 6:    6 Asset_693 2014-02-23 22:45:14 0.24515170 0.1918058     ON
```

The output looks pretty similar with the only exception being the colon after the row number. One thing to take away, though, is that it takes less time to read a `data.table` as it does a `data.frame`. This is because `data.table` only makes references to the underlying data, whereas `data.frame` copies data into a new object. This will be a recurring theme when comparing `data.table` and `data.frame`.

Another nice thing is that when you print out a large `data.table` object, you are only shown a summary and not the whole things. You wouldn't want to do the following with a huge `data.frame`:

```
DT
```

```
##           name  assetId      dateTime    signal1    signal2 signal3
##      1:         1 Asset_821 2014-08-22 22:52:14 0.73100637 0.7845882    OFF
##      2:         2 Asset_19 2014-09-18 04:22:36 0.08924429 0.2062187    OFF
##      3:         3 Asset_743 2014-09-24 18:24:37 0.56470635 0.9394910     ON
##      4:         4 Asset_969 2014-06-26 14:07:17 0.14074313 0.1588975     ON
##      5:         5 Asset_431 2014-10-06 02:19:34 0.25558412 0.3531577     ON
##      ---
## 9999996: 9999996 Asset_5 2014-07-31 19:18:11 0.42248875 2.6816182     ON
## 9999997: 9999997 Asset_711 2014-07-15 10:19:18 0.39674976 0.2533132    OFF
## 9999998: 9999998 Asset_86 2014-08-17 00:46:50 0.15199186 1.7126276    OFF
## 9999999: 9999999 Asset_482 2014-06-09 13:32:49 0.88929144 0.1391500     ON
## 10000000: 10000000 Asset_660 2014-10-16 05:51:32 0.13119666 0.7132515    OFF
```

Referencing rows

Let's now look at an example of conditionally selecting certain rows in both `data.frame` and `data.table`.

Select observations of `Asset_100` where `signal3` is 'ON':

```
system.time(h <- head(DF[DF$assetId == 'Asset_100' & DF$signal3 == 'ON', ]))
```

```
##      user  system elapsed
##      1.24    0.09     1.37
```

```
h
```

```
##           name  assetId      dateTime    signal1    signal2 signal3
## 2398      2398 Asset_100 2014-11-07 15:15:34 0.1371596 2.36482202     ON
## 4317      4317 Asset_100 2014-07-14 11:22:12 0.9412520 0.03864704     ON
## 5848      5848 Asset_100 2014-04-05 17:02:20 0.4315933 0.52792693     ON
## 7483      7483 Asset_100 2014-06-13 16:53:14 0.7889169 0.06761485     ON
## 8556      8556 Asset_100 2014-09-13 09:05:34 0.3837638 3.03475422     ON
## 11717     11717 Asset_100 2014-12-05 05:31:43 0.5391351 0.06165879     ON
```

```
system.time(h <- head(DT[assetId == 'Asset_100' & signal3 == 'ON', ]))
```

```
##      user  system elapsed
##      0.96    0.05     1.00
```

```
h
```

```
##           name  assetId      dateTime    signal1    signal2 signal3
## 1: 2398 Asset_100 2014-11-07 15:15:34 0.1371596 2.36482202     ON
## 2: 4317 Asset_100 2014-07-14 11:22:12 0.9412520 0.03864704     ON
## 3: 5848 Asset_100 2014-04-05 17:02:20 0.4315933 0.52792693     ON
## 4: 7483 Asset_100 2014-06-13 16:53:14 0.7889169 0.06761485     ON
## 5: 8556 Asset_100 2014-09-13 09:05:34 0.3837638 3.03475422     ON
## 6: 11717 Asset_100 2014-12-05 05:31:43 0.5391351 0.06165879     ON
```

Notice that within our `data.table`, we don't have to say `DT$assetId == 'Asset_100' & DT$signal3 == 'ON'`. This is because within `data.table`'s square brackets, we can reference column names directly as variables.

While in the above code `data.table` is faster than `data.frame`, it is still not the best way to subset a `data.table`. Another method involves setting a key for the `data.table` and then subsetting.

```
setkey(DT, assetId, signal3) # set key to use binary search instead of linear scan

system.time(h <- DT[list('Asset_100', 'ON'), ])
```

```
##      user      system elapsed
##         0          0         0
```

```
h
```

```
##      name  assetId      dateTime  signal1  signal2 signal3
##    1:   2398 Asset_100 2014-11-07 15:15:34 0.1371596 2.36482202      ON
##    2:   4317 Asset_100 2014-07-14 11:22:12 0.9412520 0.03864704      ON
##    3:   5848 Asset_100 2014-04-05 17:02:20 0.4315933 0.52792693      ON
##    4:   7483 Asset_100 2014-06-13 16:53:14 0.7889169 0.06761485      ON
##    5:   8556 Asset_100 2014-09-13 09:05:34 0.3837638 3.03475422      ON
##  ---
## 5038: 9984143 Asset_100 2014-08-19 00:57:41 0.2933300 0.08590503      ON
## 5039: 9985489 Asset_100 2014-03-10 12:45:33 0.1134014 1.23107313      ON
## 5040: 9986196 Asset_100 2014-09-02 06:04:58 0.5648079 0.07894776      ON
## 5041: 9991227 Asset_100 2014-06-19 06:23:32 0.5629068 0.33177273      ON
## 5042: 9998978 Asset_100 2014-03-02 10:51:36 0.8321245 0.01032277      ON
```

While setting up the key initially can take some time, all of the later subsetting is much faster. This is because `data.table` rearranges itself to allow binary search instead of linear scan. This means that instead of checking every row for these conditions, `data.table` can immediately eliminate many rows. For comparison, the computational complexity of vector scan is $O(n)$, while that of binary search is $O(\log n)$. Additionally, `data.table` performs computations by reference instead of making a copy and performing calculations on these. This is much more performant and memory efficient.

Referencing columns

Referencing columns within `data.table` is something that can seem a little confusing at first – especially when one is used to `data.frame` syntax. Take the following as an example:

```
head(DF[, 'signal1'])
```

```
## [1] 0.73100637 0.08924429 0.56470635 0.14074313 0.25558412 0.24515170
```

```
head(DT[, 'signal1'])
```

```
## [1] "signal1"
```

When we apply the same syntax from `data.frame` to `data.table`, we get something quite different. To some this may seem like a bug, but it is actually made this way by design. The second argument within `data.table`, which in `data.frame` references columns, can be an **expression** and not simply column names or indexes. So when you want to return the data from `signal1`, you can do either of the following:

```
head(DT[, signal1])
```

```
## [1] 0.3164516 0.5844961 0.3657417 0.1189339 0.6619537 0.7530080
```

```
head(DT[['signal1']])
```

```
## [1] 0.3164516 0.5844961 0.3657417 0.1189339 0.6619537 0.7530080
```

```
head(DT[, 'signal1', with=F])
```

```
##      signal1
## 1: 0.3164516
## 2: 0.5844961
## 3: 0.3657417
## 4: 0.1189339
## 5: 0.6619537
## 6: 0.7530080
```

In the first example we just referenced the column name directly since column names are treated as variables within `data.table`. In the second example, we are essentially treating `DT` as a list. In the last example, we had to say `with=F`. This is because we want to pass in a string directly to reference the column name. (*don't ask us why the argument is named "with="... this probably has come from a lengthy history*)

Why use `data.table` over `data.frame`?

Right now we have seen that `data.table` can do the same things that `data.frame` can do. And we have seen that it can do them a bit faster too. But the syntax seems weird and confusing at first. Is `data.table` really worth the extra effort?

Let's first look at the structure of `data.table`'s arguments and then look at some examples of where this structure is incredibly useful.

`data.table`'s arguments

`data.table`'s inputs - often denoted as `DT[i, j, by]` - allow the following:

- `i` allows you to evaluate conditional arguments (i.e. `signal1 > 0.5`)
- `j` allows you to select or perform expressions on columns
- `by` allows you to perform evaluations by group

`data.table`'s syntax is (in many ways) analogous to SQL. For example, you can think of the inputs to `data.table` as the following:

```
DT[where, select|update, group by][order by][...] ... [...]
```

Cool stuff in data.table

So we know that we can pass expressions to data.table's j argument. Here are a couple examples of where this could be useful:

Calculate the mean of signal1:

```
DT[, mean(signal1)]
```

```
## [1] 0.4999985
```

Calculate the mean and standard deviation of signal1:

```
DT[, list(avg=mean(signal1), sd=sd(signal1))]
```

```
##           avg           sd
## 1: 0.4999985 0.2887016
```

Create a new column called sigDif, which is the difference between signal2 and signal1:

```
head(DT[, sigDif := signal2 - signal1])
```

```
##   name assetId      dateTime  signal1  signal2 signal3  sigDif
## 1: 1652 Asset_1 2014-09-19 07:00:28 0.3164516 0.3503752812 OFF 0.0339237
## 2: 3498 Asset_1 2014-04-10 08:28:06 0.5844961 3.9385727549 OFF 3.3540767
## 3: 5981 Asset_1 2014-01-11 03:42:29 0.3657417 0.0024183812 OFF -0.3633233
## 4: 6692 Asset_1 2014-10-15 07:47:00 0.1189339 0.6390884416 OFF 0.5201546
## 5: 6832 Asset_1 2014-12-01 08:08:43 0.6619537 0.2571891844 OFF -0.4047645
## 6: 6982 Asset_1 2014-05-23 13:29:52 0.7530080 0.0008710201 OFF -0.7521370
```

Notice we use := to assign calculations to this new column. We can also use data.table's by argument to perform these calculations by group:

Calculate the mean and standard deviation of signal1 by assetId:

```
DT[, list(avg=mean(signal1), sd=sd(signal1)), by=assetId]
```

```
##           assetId      avg      sd
## 1: Asset_1 0.5022450 0.2882127
## 2: Asset_10 0.4979278 0.2920786
## 3: Asset_100 0.4992683 0.2886824
```

```
##      4: Asset_1000 0.5038555 0.2892187
##      5: Asset_101 0.4990661 0.2904023
##      ---
## 996: Asset_995 0.4956604 0.2892405
## 997: Asset_996 0.4989271 0.2886280
## 998: Asset_997 0.4967501 0.2887766
## 999: Asset_998 0.4990324 0.2888313
## 1000: Asset_999 0.4997502 0.2875915
```

Helpful links

[Introduction to the data.table package in R](#)

[FAQs about the data.table package in R](#)

[Matt Dowle's "data.table" talk at useR 2014](#)