

# MNIST Example

*Carlos Carvalho, Mladen Kolar, Rob McCulloch*

*10/21/2015*

In this example, we will explore the famous MNIST handwritten digits data set.

Data are provided by Yann LeCun and can be downloaded here: <http://yann.lecun.com/exdb/mnist/index.html>.

For convenience, the data can be also downloaded from a GitHub [repository](#). For example, you can clone the repository using the following command:

```
git clone https://github.com/ChicagoBoothML/DATA\_\_\_LeCun\_\_\_MNISTDigits.git
```

**Remark:** Data can also be downloaded from Kaggle: <https://www.kaggle.com/c/digit-recognizer/data>. Note that the data available from the Kaggle's website is not partitioned in the same way into training and test sets. Below, I will be using data from Yann LeCun's website.

The MNIST data set has a training set of 60,000 examples, and a test set of 10,000 examples. The digits have been size-normalized and centered in a fixed-size image.

Each observation is a grey-scale image sized 28 by 28 pixels. The columns are the pixel numbers, ranging from pixel 0 to pixel 783 (784 total pixels), which have elements taking values from 0 to 255 (white is 0 and 255 is black). Thus, our observations each have 784 feature values.

**The goal** is to build a model that will be presented with an image of a numerical digit (0-9) and the model must predict which digit is being shown.

## Data Import

Let us load the data.

You will need to change the code below to match the directory where you downloaded MNIST files.

```
# MNIST files in Git repository
# MNIST_DIR = "/home/mkolar/projects/mlRepos/DATA___LeCun___MNISTDigits"

# MNIST files in the current directory
MNIST_DIR = "."
```

We will also need special code to load the data set, since the data set is stored in the IDX file format. You can find more about the file format [here](#).

```
load_mnist <- function(folder) {
  load_image_file <- function(filename) {
    ret = list()
    f = file(filename, 'rb')
    readBin(f, 'integer', n=1, size=4, endian='big')
    ret$n = readBin(f, 'integer', n=1, size=4, endian='big')
    nrow = readBin(f, 'integer', n=1, size=4, endian='big')
    ncol = readBin(f, 'integer', n=1, size=4, endian='big')
```

```

    x = readBin(f, 'integer', n=ret$n*nrow*ncol, size=1, signed=F)
    ret$x = matrix(x, ncol=nrow*ncol, byrow=T)
    close(f)
    ret
  }
load_label_file <- function(filename) {
  f = file(filename, 'rb')
  readBin(f, 'integer', n=1, size=4, endian='big')
  n = readBin(f, 'integer', n=1, size=4, endian='big')
  y = readBin(f, 'integer', n=n, size=1, signed=F)
  close(f)
  y
}
train <- load_image_file(file.path(folder, 'train-images.idx3-ubyte'))
test <- load_image_file(file.path(folder, 't10k-images.idx3-ubyte'))

train$y <- load_label_file(file.path(folder, 'train-labels.idx1-ubyte'))
test$y <- load_label_file(file.path(folder, 't10k-labels.idx1-ubyte'))

list(train=train, test=test)
}

```

Using the above code, we can load the digits

```
digit.data = load_mnist(MNIST_DIR)
```

The training sample size

```
digit.data$train$n
```

```
## [1] 60000
```

Number of features

```
ncol(digit.data$train$x)
```

```
## [1] 784
```


Each image is a row of the training matrix. The following code will represent one image.

```


show_digit <- function(arr784, col=gray(12:1/12), ...) {
  image(matrix(arr784, nrow=28)[,28:1], col=col, ...)
}

show_digit(digit.data$train$x[1, ])

```



```
show_digit(digit.data$train$x[2, ])
```



Pixels are organized into images like this:

```
001 002 003 ... 026 027 028
029 030 031 ... 054 055 056
057 058 059 ... 082 083 084
|   |   |   ... |   |   |
729 730 731 ... 754 755 756
757 758 759 ... 782 783 784
```

## Logistic regression

```
library(glmnet)
if (file.exists("glmnet_mnist.RData")) {
  load("glmnet_mnist.RData")
} else {
  glm_fit = cv.glmnet(x=digit.data$train$x, y=as.factor(digit.data$train$y), family="multinomial",
                     type.logistic="modified.Newton")
  save(glm_fit, file = "glmnet_mnist.RData")
}

phat = predict(glm_fit$glmnet.fit, digit.data$test$x, s=glm_fit$lambda.1se, type = "response")
yhat = apply(phat,1,which.max)
ot = table(yhat, digit.data$test$y)
sum(diag(ot)) / 10000 # accuracy
```

```
## [1] 0.9269
```

## Random Forests

```
train = digit.data$train$x
test = digit.data$test$x
label = as.factor(digit.data$train$y)

if (file.exists("rf_mtry_28_MNIST.RData")) {
  load("rf_mtry_28_MNIST.RData")
} else {
  num_trees = 1000

  rf_28 = randomForest(
    x=train,
    y=label,
    xtests=test,
    sampsize=6000, # sample about 10% of data
    ntree=num_trees,
    mtry=28, # try 28 = sqrt(784) features at each split
    importance=TRUE,
    nodesize=100 # need this many observations in the leaf
  )

  save(rf_28,file = "rf_mtry_28_MNIST.RData")
}
rf_28
```


```
##
## Call:
## randomForest(x = train, y = label, ntree = num_trees, mtry = 28, nodesize = 100, importance = TRUE)
##               Type of random forest: classification
##               Number of trees: 1000
## No. of variables tried at each split: 28
```

```
##
##          OOB estimate of  error rate: 5.29%
## Confusion matrix:
##      0    1    2    3    4    5    6    7    8    9 class.error
## 0 5790     1   10    1    5    7   20    2   83    4 0.02245484
## 1     0 6559   63   27   15   15   10   15   31    7 0.02714328
## 2   39    9 5657   40   57    3   34   44   60   15 0.05052031
## 3   17   16  125 5648   12   89   15   59  101   49 0.07877997
## 4   15    7   21    1 5548    0   43   12   35  160 0.05032523
## 5   43   16   15   83   14 5073   63    9   69   36 0.06419480
## 6   36   14    7    0   16   62 5740    0   43    0 0.03007773
## 7    8   28   94    9   56    1    0 5895   33  141 0.05905826
## 8   17   45   43   72   37   49   31    8 5440  109 0.07024440
## 9   37   14   33   97  105   23    4   78   80 5478 0.07917297
```

Important pixels

```
varImpPlot(rf_28, type=2, n.var=20, main="Variable importance")
```

## Variable importance



Confusion matrix for the test set

```
predicted.test = predict(rf_28, test)

confusionMatrix(table(predicted.test,digit.data$test$y))
```

```

## Confusion Matrix and Statistics
##
##
## predicted.test      0      1      2      3      4      5      6      7      8      9
##      0 967      0      5      2      2      6      9      1      5      8
##      1  0 1118      0      0      1      3      3      4      1      6
##      2  0      3 978      23      2      1      1      27      7      2
##      3  0      4      9 942      0      18      0      5      7      14
##      4  0      0      11      2 929      5      6      4      6      18
##      5  3      1      1      13      1 835      7      0      5      1
##      6  4      4      7      0      6      9 925      0      9      1
##      7  1      1      9      11      0      3      0 960      4      4
##      8  5      4      10      13      7      8      7      5 912      16
##      9  0      0      2      4      34      4      0      22      18 939
##
## Overall Statistics
##
##              Accuracy : 0.9505
##              95% CI : (0.9461, 0.9547)
##      No Information Rate : 0.1135
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.945
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: 0 Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      0.9867      0.9850      0.9477      0.9327      0.9460      0.9361
## Specificity      0.9958      0.9980      0.9926      0.9937      0.9942      0.9965
## Pos Pred Value    0.9622      0.9842      0.9368      0.9429      0.9470      0.9631
## Neg Pred Value    0.9986      0.9981      0.9940      0.9924      0.9941      0.9938
## Prevalence        0.0980      0.1135      0.1032      0.1010      0.0982      0.0892
## Detection Rate    0.0967      0.1118      0.0978      0.0942      0.0929      0.0835
## Detection Prevalence 0.1005      0.1136      0.1044      0.0999      0.0981      0.0867
## Balanced Accuracy 0.9913      0.9915      0.9702      0.9632      0.9701      0.9663
##
##              Class: 6 Class: 7 Class: 8 Class: 9
## Sensitivity      0.9656      0.9339      0.9363      0.9306
## Specificity      0.9956      0.9963      0.9917      0.9907
## Pos Pred Value    0.9585      0.9668      0.9240      0.9179
## Neg Pred Value    0.9963      0.9925      0.9931      0.9922
## Prevalence        0.0958      0.1028      0.0974      0.1009
## Detection Rate    0.0925      0.0960      0.0912      0.0939
## Detection Prevalence 0.0965      0.0993      0.0987      0.1023
## Balanced Accuracy 0.9806      0.9651      0.9640      0.9606

```