

# Trees

Carlos Carvalho, Mladen Kolar and Rob McCulloch

1. Trees
2. Regression Trees
3. Classification Trees
4. Trees: A Summary

# 1. Trees

Tree based methods are a major player in data-mining.

Good:

- ▶ flexible fitters, capture non-linearity and interactions.
- ▶ do not have to think about scale of  $x$  variables.
- ▶ handles categorical and numeric  $y$  and  $x$  very nicely.
- ▶ fast.
- ▶ interpretable (when small).

Bad:

Not the best in out-of-sample predictive performance  
(*but not bad!!*).

*But,*

If we **bag** or **boost** trees, we can get the best off-the-shelf prediction available.

Bagging and Boosting are *ensemble methods* that combine the fit from many (hundreds, thousands) of tree models to get an overall predictor.

## 2. Regression Trees



Let's look at a simple 1-dimensional example so that we can see what is going on.

We'll use the Boston housing data and relate  $x=\text{lstat}$  to  $y=\text{medval}$ .

At left is the *tree* fit to the data.



At each *interior node* there is a decision rule of the form  $\{x < c\}$ . If  $x < c$  you go left, otherwise you go right.

Each observation is sent down the tree until it hits a bottom node or *leaf* of the tree.



The set of bottom nodes gives us a partition of the predictor ( $x$ ) space into disjoint regions. At right, the vertical lines display the partition. With just one  $x$ , this is just a set of intervals.

Within each region (interval) we compute the average of the  $y$  values for the subset of training data in the region. This gives us the step function which is our  $\hat{f}$ . The  $\bar{y}$  values are also printed at the bottom nodes (left plot).





To predict, we just use our step function estimate of  $f(x)$ .

Equivalently, we drop  $x$  down the tree until it lands in a leaf and then predict the average of the  $y$  values for the training observations in the same leaf.

## A Tree with Two Explanatory Variables

Here is a tree with  $x = (x_1, x_2) = (\text{lstat}, \text{dis})$  and  $y = \text{medv}$ .

Now the decision rules can use either of the two  $x$ 's.



At right is the *partition* of the  $x$  space corresponding to the set of bottom nodes (leaves).

The average  $y$  for training observations assigned to a region is printed in each region and at the bottom nodes.




This is the regression function given by the tree.

It is a step function which can seem dumb, but it delivers non-linearity *and* interactions in a simple way and works with a lot of variables.



Notice the interaction.

The effect of `dis` depends on `lstat`!!




# The California Housing Data

Here is a tree with 50 bottom nodes fit to the California Housing data using only longitude and latitude.



Don't extrapolate into the ocean!

Here is a view of the fit using the map of the state.  
(units are dollars, the logMedVal was exponentiated for the labels).



## In R:

```
#-----  
#load tree package (and MASS), attach Boston data  
library(tree)  
library(MASS)  
#data(Boston) #don't need this  
attach(Boston)  
  
#-----  
#fit a tree to boston data just using lstat.  
#first get a big tree using a small value of mindev  
temp = tree(medv~lstat,data=Boston,mindev=.0001)  
cat("first big tree size: \n")  
print(length(unique(temp$where)))  
#if the tree is too small, make mindev smaller!!  
  
#-----  
#then prune it down to one with 7 leaves  
boston.tree=prune.tree(temp,best=7)  
cat("pruned tree size: \n")  
print(length(unique(boston.tree$where)))
```

```

#-----
#plot the tree
plot(boston.tree,type="uniform")
text(boston.tree,col="blue",label=c("yval"),cex=.8)

#-----
#plot data with fit
#get fit
boston.fit = predict(boston.tree) #get training fitted values
#plot fit
plot(lstat,medv,cex=.5,pch=16) #plot data
oo=order(lstat)
lines(lstat[oo],boston.fit[oo],col="red",lwd=3) #step function fit

#-----
#predict at lstat = 15 and 25.
preddf = data.frame(lstat=c(15,25))
yhat = predict(boston.tree,preddf)
points(preddf$lstat,yhat,col="blue",pch="*",cex=3)

```

Let's fit the tree using lstat and dis.

```
#-----  
df2=Boston[,c(8,13,14)] #pick off dis,lstat,medv  
print(names(df2))  
  
#-----  
#big tree  
temp = tree(medv~.,df2,mindev=.0001)  
cat("first big tree size: \n")  
print(length(unique(temp$where)))  
  
#-----  
#then prune it down to one with 7 leaves  
boston.tree=prune.tree(temp,best=7)  
cat("pruned tree size: \n")  
print(length(unique(boston.tree$where)))  
  
#-----  
# plot tree and partition in x.  
par(mfrow=c(1,2))  
#plot tree  
plot(boston.tree,type="u")  
text(boston.tree,col="blue",label=c("yval"),cex=.8)  
#plot 2-dimesional partition in (x1,x2) = (lstat,dis)  
partition.tree(boston.tree)
```

```

#-----
#let's compare in-sample fits from our two trees with each other and y
boston.fit2 = predict(boston.tree)
fmat = cbind(medv,boston.fit,boston.fit2)
colnames(fmat)=c("y=medv","treel","treeld")
pairs(fmat)
print(cor(fmat))

#-----
#predict at lstat = 15 and 25 and dis = 2 both times
preddf=data.frame(lstat=c(15,25),dis=c(2,2))
yhat2 = predict(boston.tree,preddf)
cat("predictions are:\n")
print(yhat2)

```

Lets try  $p = 4$  with nox, rm, ptratio, and lstat.

```
#-----  
df4=Boston[,c(5,6,11,13,14)] #pick off variables  
print(names(df4))  
temp = tree(medv~.,df4,mindev=.0001)  
cat("first big tree size: \n")  
print(length(unique(temp$where)))  
  
#-----  
#then prune it down to one with 15 leaves (picked 15 arbitrarily)  
boston.tree4=prune.tree(temp,best=15)  
cat("pruned tree size: \n")  
print(length(unique(boston.tree4$where)))  
  
#-----  
#plot tree  
par(mfrow=c(1,1))  
plot(boston.tree4,type="u")  
text(boston.tree4,col="blue",label=c("yval"),cex=.8)  
  
#-----  
#compare fits  
fmat4=cbind(fmat,predict(boston.tree4))  
colnames(fmat4)[4]="tree4"  
pairs(fmat4)  
print(cor(fmat4))
```



### 3. Classification Trees


Let's do a tree for a classification problem.

We'll use the hockey penalty data.

The response is whether or not the next penalty is on the other team and  $x$  is a bunch of stuff about the game situation (the score ...).

In addition, this time some of our predictors (features,  $x$ 's) are categorical.

Here is the tree:



- ▶ Each bottom node gives the fraction of training data in the two outcome categories. Think of it as  $\hat{p}$  for the kind of  $x$  associated with that bottom node.
- ▶ The form of the decision rule can't be  $x < c$  for categorical variables. We pick a subset of the levels to go left. `inrow2:0` means all the observations with `inrow2` in the category labeled 0 go left.

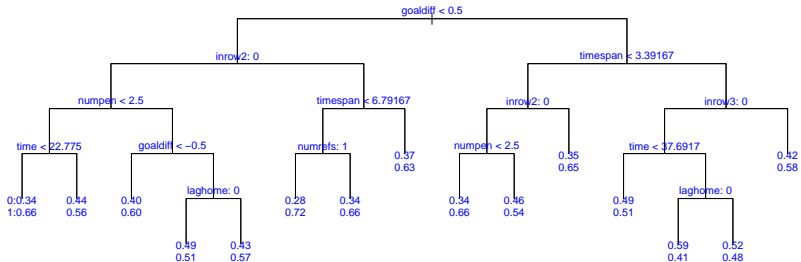
There is a lot of fit!!!

*if:*

- ▶ if you are not winning
- ▶ you had the last two penalties
- ▶ it has not been long since the last call
- ▶ and there is only 1 referee

*then:*

there is a 72% chance the next call will be on the other team.



Whilst there is another game situation where the chance the next call is on the other team is only 41%.

## 4. Trees: A Summary

### Trees:

- ▶ Trees use recursive binary splits to partition the predictor space.
- ▶ Each binary split consists of a decision rule which sends  $x$  left or right.
- ▶ For numeric  $x_i$ , the decision rule is of the form if  $x_i < c$ .
- ▶ For categorical  $x_i$ , the rule lists the set of categories sent left.
- ▶ The set of bottom nodes (or leaves) give a partition of the  $x$  space.
- ▶ To predict, we drop an out-of-sample  $x$  down the tree until it lands in a bottom node.
- ▶ For numeric  $y$ , we predict the average  $y$  value for the training data that ended up in the bottom node.
- ▶ For categorical  $y$  we use the category proportions for the training data that ended up in the bottom node.

## Good:

- ▶ Handles categorical/numeric  $x$  and  $y$  nicely.
- ▶ Don't have to think about the scale of  $x$ 's !!!
- ▶ Computationally fast ("scales").
- ▶ Small trees are interpretable.
- ▶ Variable selection.

## Bad:

- ▶ Step function is crude, does not give the best predictive performance.
- ▶ Hard to assess uncertainty.
- ▶ Big trees are not interpretable.